
THE POWER OF INTELLIGENT FLOWS

REAL-TIME IOT BOTNET CLASSIFICATION
WITH APACHE NIFI

Andre Fucs de Miranda - Fluenda
Andy LoPresto - Hortonworks

Agenda

- Who are the two blokes in front of you
 - A brief prologue
 - Logs! Logs! Logs!
 - The challenge
 - The solution
 - Wrapping up
-

Who are the two blokes in front of you

Andre Fucs de Miranda

- Nearly 20 years working with information cyber security

- Logging aficionado (i.e. security data engineer)

- Apache NiFi PMC Member



@trixpan



@trixpan

Andy LoPresto

- Financial security & device firmware at Apple, TigerText, etc.

- PII, PCI & EPHI encryption & cracking

- Apache NiFi PMC Member



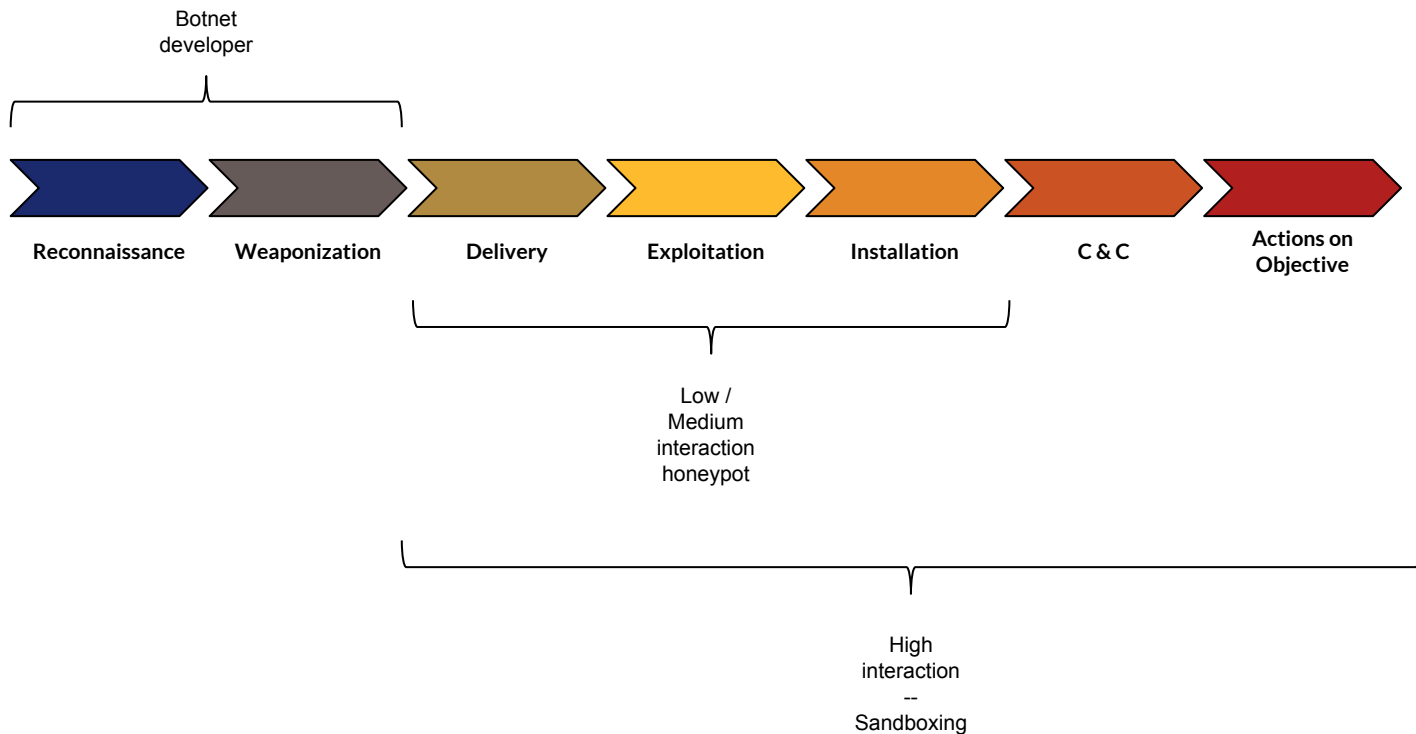
@yolohey



@alopresto

A brief prologue

The Botnet Kill Chain & the Honeypot



The Botnet Kill Chain & the Honeytrap



Delivery

Exploitation

Installation

Step 1

Logon to
system

Step 2


Execute
predefined
sequence of
commands

Step 3

Try to install
some sort of
persistence


The demo environment

- A handful of EC2 instances running:
 - Cowrie - Medium interaction SSH / Telnet honeypot
 - MiNiFi
- An EC2 instance running:
 - NiFi 1.3.0 (with security enabled)



Flow Design Approach

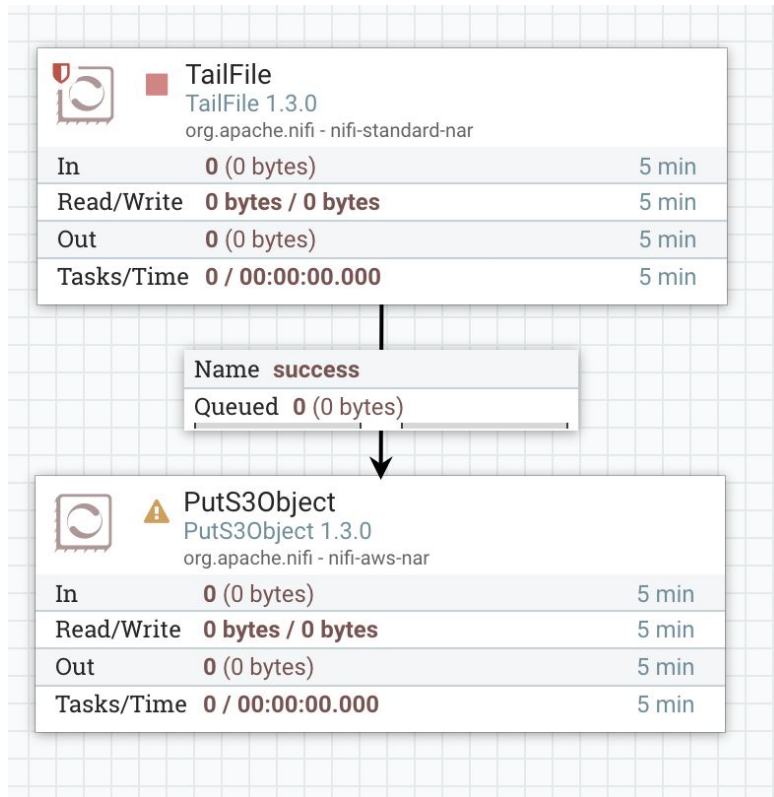
- Don't be prescriptive
- Treat everything as data
- Don't be limited by prior expectations
- *Start from the end*



Logs! Logs! Logs!

MiNiFi Process Group

- Tailing a log file being written by cowrie
- Pushing to Amazon S3
 - Could stream via NiFi Site to Site
 - MiNiFi extensibility
 - Shows multiple capabilities
 - Decoupled/no lock in



The data being ingested

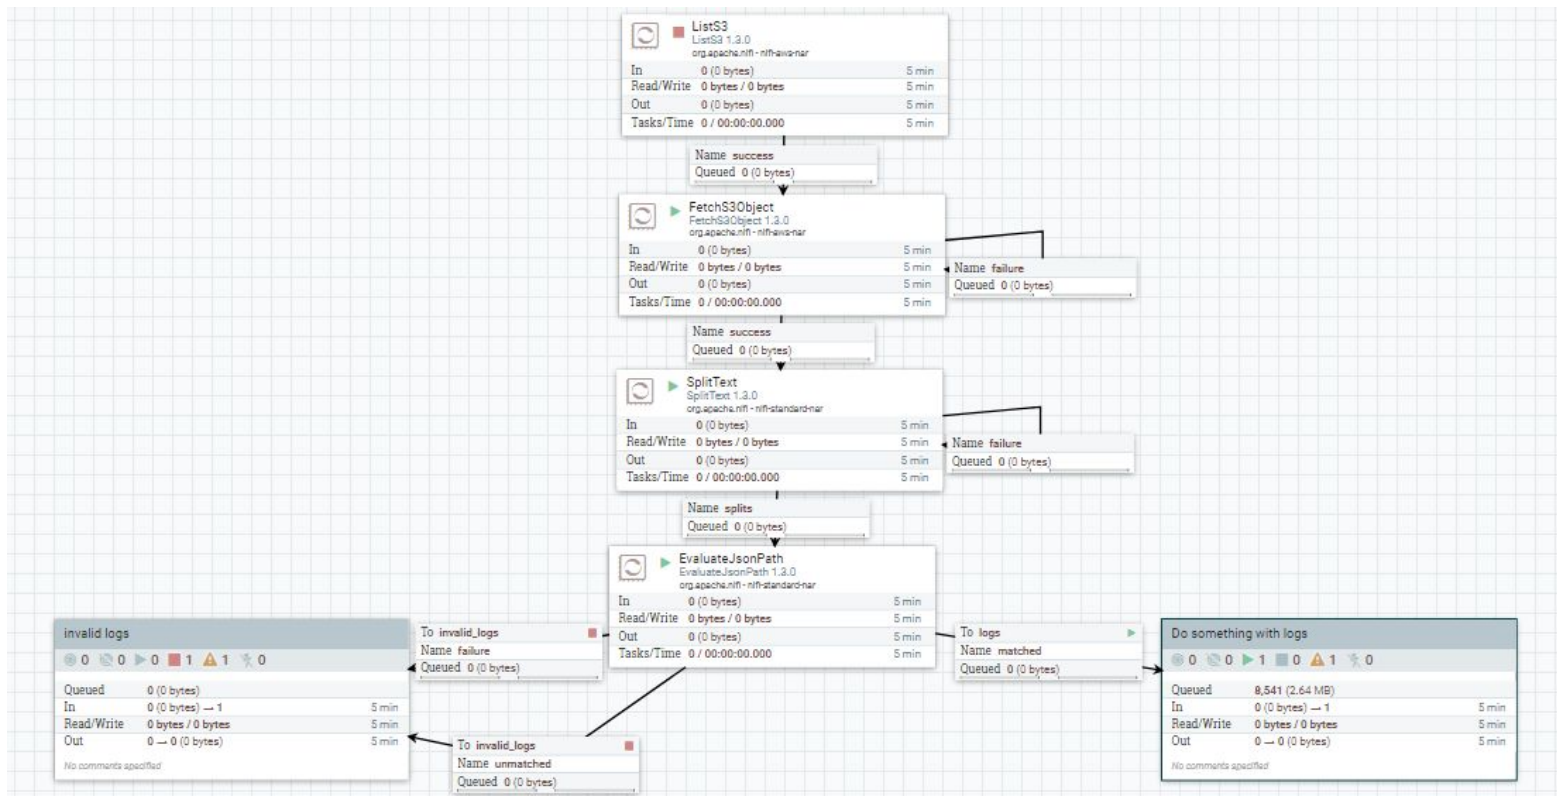
- Cowrie logs include:
 - Username / Password
 - Commands executed (and parameters)
 - Files downloaded
- Single line JSON entries
 - Easy to parse
 - Textbook machine readable log format
 - Perfect match to NiFi processors such as:
 - SplitText
 - EvaluateJSONPath

Cowrie log example

```
{
  "eventid": "cowrie.log.closed", "timestamp": "2017-09-20T00:37:24.713248Z", "message": "Closing TTY Log: log/tty/20170920-003522-None-2128i.log after 122 se",
  "eventid": "cowrie.session.closed", "timestamp": "2017-09-20T00:37:24.719938Z", "message": "Connection lost after 127 seconds", "system": "CowrieTelnetTrans",
  "eventid": "cowrie.session.connect", "timestamp": "2017-09-20T00:35:17.055727Z", "session": "82f705bca22c", "message": "New connection: 51.15.221.208:42242",
  "eventid": "cowrie.login.success", "username": "root", "timestamp": "2017-09-20T00:35:21.883371Z", "message": "login attempt [root/1234] succeeded", "system": "CowrieTelnetTransport",
  "eventid": "cowrie.log.open", "timestamp": "2017-09-20T00:35:22.209562Z", "message": "Opening TTY Log: log/tty/20170920-003522-None-2128i.log", "ttylog": "1",
  "eventid": "cowrie.command.input", "timestamp": "2017-09-20T00:35:22.290251Z", "message": "CMD: enable", "system": "CowrieTelnetTransport,2128,51.15.221.208",
  "eventid": "cowrie.command.success", "timestamp": "2017-09-20T00:35:22.291474Z", "message": "Command found: enable ", "system": "CowrieTelnetTransport,2128,51.15.221.208",
  "eventid": "cowrie.command.input", "timestamp": "2017-09-20T00:35:22.368973Z", "message": "CMD: shell", "system": "CowrieTelnetTransport,2128,51.15.221.208",
  "eventid": "cowrie.command.failed", "timestamp": "2017-09-20T00:35:22.370123Z", "message": "Command not found: shell", "system": "CowrieTelnetTransport,2128,51.15.221.208",
  "eventid": "cowrie.command.input", "timestamp": "2017-09-20T00:35:22.370877Z", "message": "CMD: sh", "system": "CowrieTelnetTransport,2128,51.15.221.208",
  "eventid": "cowrie.command.success", "timestamp": "2017-09-20T00:35:22.371830Z", "message": "Command found: sh ", "system": "CowrieTelnetTransport,2128,51.15.221.208",
  "eventid": "cowrie.command.input", "timestamp": "2017-09-20T00:35:22.485810Z", "message": "CMD: /bin/busybox ECCHI", "system": "CowrieTelnetTransport,2128,51.15.221.208",
  "eventid": "cowrie.command.success", "timestamp": "2017-09-20T00:35:22.486541Z", "message": "Command found: /bin/busybox ECCHI", "system": "CowrieTelnetTransport,2128,51.15.221.208",
  "eventid": "cowrie.command.input", "timestamp": "2017-09-20T00:35:22.568321Z", "message": "CMD: /bin/busybox ps; /bin/busybox ECCHI", "system": "CowrieTelnetTransport,2128,51.15.221.208",
  "eventid": "cowrie.command.success", "timestamp": "2017-09-20T00:35:22.569059Z", "message": "Command found: /bin/busybox ps", "system": "CowrieTelnetTransport,2128,51.15.221.208",
  "eventid": "cowrie.command.success", "timestamp": "2017-09-20T00:35:22.570106Z", "message": "Command found: ps", "system": "CowrieTelnetTransport,2128,51.15.221.208",
  "eventid": "cowrie.command.success", "timestamp": "2017-09-20T00:35:22.571035Z", "message": "Command found: /bin/busybox ECCHI", "system": "CowrieTelnetTransport,2128,51.15.221.208",
  "eventid": "cowrie.command.input", "timestamp": "2017-09-20T00:35:22.655910Z", "message": "CMD: /bin/busybox cat /proc/mounts; /bin/busybox ECCHI", "system": "CowrieTelnetTransport,2128,51.15.221.208",
  "eventid": "cowrie.command.success", "timestamp": "2017-09-20T00:35:22.656838Z", "message": "Command found: /bin/busybox cat /proc/mounts", "system": "CowrieTelnetTransport,2128,51.15.221.208",
  "eventid": "cowrie.command.success", "timestamp": "2017-09-20T00:35:22.657870Z", "message": "Command found: cat /proc/mounts", "system": "CowrieTelnetTransport,2128,51.15.221.208",
  "eventid": "cowrie.command.success", "timestamp": "2017-09-20T00:35:22.658683Z", "message": "Command found: /bin/busybox ECCHI", "system": "CowrieTelnetTransport,2128,51.15.221.208",
  "eventid": "cowrie.command.input", "timestamp": "2017-09-20T00:35:22.746314Z", "message": "CMD: /bin/busybox echo -e '\\x6b\\x61\\x6d\\x69' > /.nippon; /bin"}

```

Simple Cowrie log ingestion with NiFi



The challenge

The challenge



The challenge

- Logs in isolation rarely will provide the reader with a meaningful view over what is happening
 - Verbosity means sensors generate lots of “events”, but who cares about a bot trying to ``cat /proc/mounts`` ?
 - Bots use semi-random values to make detection more difficult.
-

The solution

Logs are data too...

View as: formatted

```
113 "sensor": "a927e8b28666"
114 }, {
115   "eventid": "cowrie.command.input",
116   "timestamp": "2017-09-18T11:45:25.025835Z",
117   "message": "CMD: cat /proc/mounts; /bin/busybox LSUCT",
118   "system": "CowrieTelnetTransport,787,121.237.129.163",
119   "isError": 0,
120   "src_ip": "121.237.129.163",
121   "session": "21caf72c6358",
122   "input": "cat /proc/mounts; /bin/busybox LSUCT",
123   "sensor": "a927e8b28666"
124 }, {
125   "eventid": "cowrie.command.success",
126   "timestamp": "2017-09-18T11:45:25.027153Z",
127   "message": "Command found: cat /proc/mounts",
128   "system": "CowrieTelnetTransport,787,121.237.129.163",
129   "isError": 0,
130   "src_ip": "121.237.129.163",
131   "session": "21caf72c6358",
132   "input": "cat /proc/mounts",
133   "sensor": "a927e8b28666"
134 }, {
135   "eventid": "cowrie.command.success",
136   "timestamp": "2017-09-18T11:45:25.028091Z",
137   "message": "Command found: /bin/busybox LSUCT",
138   "system": "CowrieTelnetTransport,787,121.237.129.163",
139   "isError": 0,
140   "src_ip": "121.237.129.163",
141   "session": "21caf72c6358",
142   "input": "/bin/busybox LSUCT",
143   "sensor": "a927e8b28666"
144 }, {
145   "eventid": "cowrie.command.input",
146   "timestamp": "2017-09-18T11:45:25.367150Z",
147   "message": "CMD: cd /dev/shm; cat .s || cp /bin/echo .s; /bin/busybox LSUCT",
148   "system": "CowrieTelnetTransport,787,121.237.129.163",
149   "isError": 0,
150   "src_ip": "121.237.129.163",
151   "session": "21caf72c6358",
152   "input": "cd /dev/shm; cat .s || cp /bin/echo .s; /bin/busybox LSUCT",
153   "sensor": "a927e8b28666"
154 }, {
155   "eventid": "cowrie.command.success",
```



View as: formatted


```
147 }, {
148   "eventid": "cowrie.command.input",
149   "timestamp": "2017-09-17T04:06:39.670673Z",
150   "message": "CMD: cat /proc/mounts; /bin/busybox XUSRH",
151   "system": "CowrieTelnetTransport,93,94.51.110.74",
152   "isError": 0,
153   "src_ip": "94.51.110.74",
154   "session": "4c047bbc016c",
155   "input": "cat /proc/mounts; /bin/busybox XUSRH",
156   "sensor": "a927e8b28666"
157 }, {
158   "eventid": "cowrie.command.success",
159   "timestamp": "2017-09-17T04:06:39.672190Z",
160   "message": "Command found: cat /proc/mounts",
161   "system": "CowrieTelnetTransport,93,94.51.110.74",
162   "isError": 0,
163   "src_ip": "94.51.110.74",
164   "session": "4c047bbc016c",
165   "input": "cat /proc/mounts",
166   "sensor": "a927e8b28666"
167 }, {
168   "eventid": "cowrie.command.success",
169   "timestamp": "2017-09-17T04:06:39.673206Z",
170   "message": "Command found: /bin/busybox XUSRH",
171   "system": "CowrieTelnetTransport,93,94.51.110.74",
172   "isError": 0,
173   "src_ip": "94.51.110.74",
174   "session": "4c047bbc016c",
175   "input": "/bin/busybox XUSRH",
176   "sensor": "a927e8b28666"
177 }, {
178   "eventid": "cowrie.command.input",
179   "timestamp": "2017-09-17T04:06:39.859611Z",
180   "message": "CMD: cd /dev/shm; cat .s || cp /bin/echo .s; /bin/busybox XUSRH",
181   "system": "CowrieTelnetTransport,93,94.51.110.74",
182   "isError": 0,
183   "src_ip": "94.51.110.74",
184   "session": "4c047bbc016c",
185   "input": "cd /dev/shm; cat .s || cp /bin/echo .s; /bin/busybox XUSRH",
186   "sensor": "a927e8b28666"
187 }, {
188   "eventid": "cowrie.command.success",
189   "timestamp": "2017-09-17T04:06:39.860568Z",
```



This looks familiar...

Bill-36641

Spam x

 Rosa Ogram <RosaOgram@cridersvillehealthcare.com>
to conductorn

Aug 25 ☆  

Why is this message in Spam? It's similar to messages that were detected by our spam filters. [Learn more](#)



Images are not displayed. [Display images below](#)


Hi,
[Here](#) is a copy of your bill.
Thank you & have a great weekend!
...



Rosa Ogram

Bill-85882

Spam x

 Gretchen Jagger <GretchenJagger@pineiroproductions.com>
to als

Aug 25 ☆  

Why is this message in Spam? It contains content that's typically used in spam messages. [Learn more](#)

Images are not displayed. [Display images below](#)

Hi,
[Here](#) is a copy of your bill.
Thank you & have a great weekend!
...

Gretchen Jagger

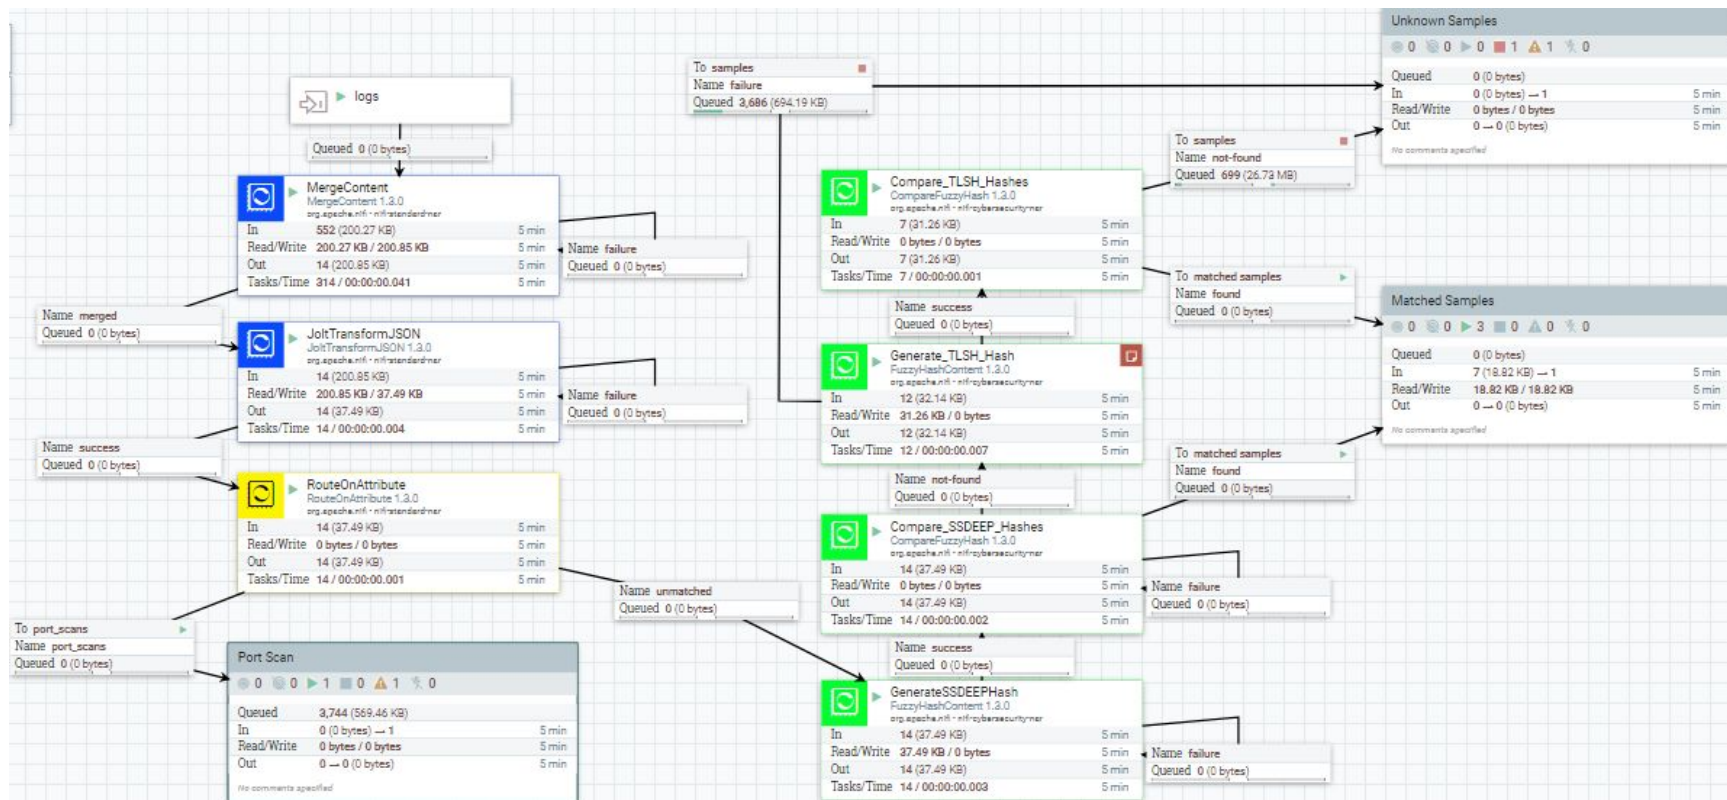
Locality-sensitive hashing

A type of algorithm that can be used to “group” similar items together and may provide a similarity score between two particular items.

Areas of application:

- Genome-wide association study
 - Anti-spam (e.g. TLSH, Spamsum/SSDeep)
 - Near-duplicate detection
 - etc
-

NiFi + SpamSum + TLSH = WIN!



Wrapping up

Key points

- Treat everything as data
- Be flexible on how you build your data flows.
- Apparently unrelated domains may speed up your results
- Use MiNiFi to aggregate data at the edge whenever possible
- NiFi rocks!*

* Disclaimer: We may be a bit biased...

Future Steps

- Automate IP blocking & firewall rules (ML)
 - Continuously update signature definition list with new sigs
 - Analyze epidemiology & spread vectors
 - Follow evolution of malware families
 - Support attribution of samples
-

Further reading

Mysterious Hajime botnet has pwned 300,000 IoT devices

https://www.theregister.co.uk/2017/04/27/hajime_iot_botnet/

Identifying unknown files by using fuzzy hashing

<https://www.honeynet.org/node/811>

Classifying Malware using Import API and Fuzzy Hashing – impfuzzy

<http://blog.jpcert.or.jp/2016/05/classifying-mal-a988.html>

Template and samples:

https://github.com/fluenda/dataworks_summit_iot_botnet

Thank you
